



UNIVERSITY OF  
BIRMINGHAM

# ARTIFICIAL INTELLIGENCE 1 OPTIMISATION

DR LEONARDO STELLA

SCHOOL OF COMPUTER SCIENCE

2023/2024 - Week 9

# AIMS OF THE SESSION

This session aims to help you:

- Understand how to formulate an optimisation problem
- Explain the steps involved in Hill Climbing
- Compare the performance of search and optimisation algorithms

# OUTLINE

1 Optimisation

2 Hill Climbing

# OUTLINE

1 Optimisation

2 Hill Climbing

# OPTIMISATION

- Until now, we have addressed problems in fully observable, deterministic, static, known environments
- In this lecture, we will relax some of these constraints and look at optimisation problems
- **Optimisation** problems arise in many disciplines, including computer science, engineering, operations research, etc.
- In an optimisation problem, one seeks to find a solution that minimises or maximises an **objective function** (or more than one)
- Depending on the problem, there can be some **constraints** that must be satisfied for a solution to be **feasible**

# FORMULATING AN OPTIMISATION PROBLEM

- The canonical representation of an optimisation problem is like in the following:

$$\begin{aligned} \min / \max \quad & f(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n, \end{aligned}$$

- where  $f(x)$  the objective function,  $x$  is the design variable
- $g_i(x)$  and  $h_j(x)$  are the constraints
- The **search space** of the problem is the space of all possible  $x$  values

# FORMULATING AN OPTIMISATION PROBLEM

- The **design variables** represent a candidate solution
- Therefore, the **search space** of a candidate solution is defined by the design variables
- The **objective function** defines the cost (or quality) of a solution
- This function is the one to be optimised (through min or max)
- The solution must optionally satisfy a number of **constraints**
- They define the feasibility of the solution

# OUTLINE

1 Optimisation

2 Hill Climbing



# LOCAL SEARCH

- We can now formalise an optimisation problem
- In a search problem, we want to find paths through the search space
- For example, finding the shortest path between Arad and Bucharest
- In the previous slides, we formalised the corresponding optimisation problem (or one possible way)
- However, sometimes we are interested only in the final state (e.g., the 8-queen problem discussed in this week's tutorial session)


# LOCAL SEARCH

- **Local search algorithms** (also referred to as **optimisation algorithms**) operate by searching from an initial (randomised) state to neighbouring states
- These algorithms **do not** keep track of the paths nor the states that have been reached/visited
- This means they are not systematic, but they have advantages:
  - they use very **little memory**
  - they can often find **reasonable solutions** in large or infinite state spaces
- Important applications include: integrated-circuit design, factory floor layout, job shop scheduling, portfolio management, etc.

# LOCAL SEARCH

**Example.** Consider the problem of maximising an objective function whose value is defined as the elevation on the y-axis in the figure below (against the state space on the x-axis). The aim is to find the highest peak, i.e., a global maximum. This process is called **hill climbing**.

**Example.** Consider the opposite problem. The objective function now represents a cost that we want to minimise. The aim is to find the lowest valley, i.e., a global minimum. This process is called **gradient descent**.



# HILL CLIMBING

- **Hill Climbing** is one of the most popular optimisation algorithms:
  - Generate initial candidate solution at random
    - Generate neighbour solutions and move to the one with highest value
    - If no neighbour has a higher value of the current solution, terminate
    - Otherwise, repeat with a new best neighbour solution
- Hill climbing does not look beyond the immediate neighbours of the current state, making it a **greedy algorithm**
- This is equivalent “to trying to find the top of Mount Everest in a thick fog while suffering from amnesia”

# HILL CLIMBING

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  problem.INITIAL  
  while true do  
    neighbor  $\leftarrow$  a highest-valued successor state of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
    current  $\leftarrow$  neighbor
```

# HILL CLIMBING

- To be able to apply hill climbing, we need to design the following components of the algorithm:
  - **Representation**, how to store the design variable, e.g., boolean, integer, float, array  
A good representation should facilitate the application of the initialisation procedure and neighbourhood operator
  - **Initialisation procedure**, how to pick an initial solution  
Usually, this involves selecting one at random
  - **Neighbourhood operator**, how to generate neighbour solutions

# PERFORMANCE OF HILL CLIMBING

- Let us evaluate the performance of hill climbing.
  - **Completeness:** hill climbing **is not complete**, as it depends on the problem formulation and design of the algorithm
  - **Optimality:** hill climbing **is not optimal** (can get stuck in a local maximum/optimum)
  - **Time complexity:**  $\mathcal{O}(mnp)$ , where  $m$  is the maximum number of iterations,  $n$  is the maximum number of neighbours, each of which takes  $\mathcal{O}(p)$  to generate
  - **Space complexity:**  $\mathcal{O}(nq + r) = \mathcal{O}(nq)$ , where the variable takes  $\mathcal{O}(q)$  and  $r$  represents the space to generate the neighbours sequentially (which is negligible compared to  $n$  and  $q$ )

# HILL CLIMBING

- Pros:

- Can **rapidly** find a good solution by improving over a bad initial state (greedy)
- **Low time and space complexity** compared to search algorithms
- **Does not require** problem-specific heuristics
- Start from **a candidate solution**, instead of building it step-by-step

- Cons:

- **Not guaranteed** to be complete, nor optimal
- Can **get stuck** in local maxima and plateaus



# References



Russell, A. S., and Norvig, P., *Artificial Intelligence A Modern Approach*, 4<sup>th</sup> Edition. Prentice Hall.



Chapter 4 – “Search in Complex Environments”, Section 4.1 “Local Search and Optimization Problems” up to and including Section 4.1.1 “Hill Climbing Search”.

# AIMS OF THE SESSION

You should now be able to:

- Understand how to formulate an optimisation problem
- Explain the steps involved in Hill Climbing
- Compare the performance of search and optimisation algorithms



Thank you!