



UNIVERSITY OF
BIRMINGHAM

ARTIFICIAL INTELLIGENCE 1
OPTIMISATION AND CONSTRAINTS
HANDLING

DR LEONARDO STELLA

SCHOOL OF COMPUTER SCIENCE

2023/2024 - Week 10

AIMS OF THE SESSION

This session aims to help you:

- Understand the weaknesses in Hill Climbing and how to mitigate them
- Explain the steps involved in Simulated Annealing
- Deal with constraints in optimisation problems

OUTLINE

- 1 Recap: Informed Search and Optimisation
- 2 Hill Climbing: Variants
- 3 Simulated Annealing
- 4 Constraint Handling

OUTLINE

- 1 Recap: Informed Search and Optimisation
- 2 Hill Climbing: Variants
- 3 Simulated Annealing
- 4 Constraint Handling

RECAP: QUIZ

Q1. Which evaluation function $f(n)$ does \mathbf{A}^* use?

- ☐ $f(n) = c(n)$, where $c(n)$ is the cost to get to n
- ☐ $f(n) = h(n)$, where $h(n)$ is the heuristic
- ☐ $f(n) = h(n) - c(n)$, where $h(n)$ is the heuristic and $c(n)$ is the cost to get to n
- ☐ $f(n) = h(n)^2$, where $h(n)$ is the heuristic
- ☐ None of the above

RECAP: QUIZ


Q1. Which evaluation function $f(n)$ does \mathbf{A}^* use?

- ☐ $f(n) = c(n)$, where $c(n)$ is the cost to get to n
- ☐ $f(n) = h(n)$, where $h(n)$ is the heuristic
- ☐ $f(n) = h(n) - c(n)$, where $h(n)$ is the heuristic and $c(n)$ is the cost to get to n
- ☐ $f(n) = h(n)^2$, where $h(n)$ is the heuristic
- ☒ None of the above

RECAP: QUIZ

Q2. Given the heuristic in the table below and the cost to reach each city on the map (initial state: Sibiu), which one is the correct value of $f(\text{Pitesti})$ for **greedy best-first search**?

- ☐ $f(\text{Pitesti}) = 100$
- ☐ $f(\text{Pitesti}) = 197$
- ☐ $f(\text{Pitesti}) = 177$
- ☐ $f(\text{Pitesti}) = 277$
- ☐ None of the above




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

RECAP: QUIZ

Q2. Given the heuristic in the table below and the cost to reach each city on the map (initial state: Sibiu), which one is the correct value of $f(\text{Pitesti})$ for **greedy best-first search**?

- ☒ $f(\text{Pitesti}) = 100$
- ☐ $f(\text{Pitesti}) = 197$
- ☐ $f(\text{Pitesti}) = 177$
- ☐ $f(\text{Pitesti}) = 277$
- ☐ None of the above




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

RECAP: QUIZ

Q3. Given the heuristic in the table below and the cost to reach each city on the map (initial state: Sibiu), which one is the correct value of $f(\text{Pitesti})$ for A^* ?

- ☐ $f(\text{Pitesti}) = 100$
- ☐ $f(\text{Pitesti}) = 197$
- ☐ $f(\text{Pitesti}) = 177$
- ☐ $f(\text{Pitesti}) = 277$
- ☐ None of the above




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

RECAP: QUIZ

Q3. Given the heuristic in the table below and the cost to reach each city on the map (initial state: Sibiu), which one is the correct value of $f(\text{Pitesti})$ for A^* ?

- ☐ $f(\text{Pitesti}) = 100$
- ☐ $f(\text{Pitesti}) = 197$
- ☐ $f(\text{Pitesti}) = 177$
- ☒ $f(\text{Pitesti}) = 277$
- ☐ None of the above



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

RECAP: QUIZ

Q4. What are the components of an optimisation problem formulation?

- ☐ Design variables
- ☐ State space
- ☐ Objective functions
- ☐ Constraints
- ☐ None of the above

RECAP: QUIZ

Q4. What are the components of an optimisation problem formulation?

- ☒ Design variables
- ☐ State space
- ☒ Objective functions
- ☒ Constraints
- ☐ None of the above

RECAP: QUIZ

- Q5. Local search algorithms **do not** keep track of the paths nor the states that have been reached/visited.
- ☐ True
 - ☐ False

RECAP: QUIZ

- Q5. Local search algorithms **do not** keep track of the paths nor the states that have been reached/visited.
- ☒ True
 - ☐ False

RECAP: QUIZ

Q6. Hill climbing is a greedy algorithm.

- ☐ True
- ☐ False

RECAP: QUIZ

Q6. Hill climbing is a greedy algorithm.

- ☒ True
- ☐ False

OUTLINE

- 1 Recap: Informed Search and Optimisation
- 2 Hill Climbing: Variants**
- 3 Simulated Annealing
- 4 Constraint Handling

WEAKNESSES OF HILL CLIMBING

- Last week, we discussed pros and cons of hill climbing
- Can we mitigate some of its weaknesses?
 - **W**: it can run for ever if the problem is not properly formulated
 - **S**: we run the algorithm for a maximum number of iterations m
 - Indeed, we used m to determine its time complexity: $\mathcal{O}(mnp)$
 - **W**: not guaranteed to be complete, nor optimal
 - **S**: as long as we are satisfied with a good solution, it can find it quickly
 - **W**: it can get stuck in local maxima and plateaus
 - **S**: there are variants of the algorithm to deal with this

HILL CLIMBING: VARIANTS

- Let us consider an n -dimensional problem
- The gradient determines the uphill moves along the n dimensions
- **Stochastic Hill Climbing**
 - This variant chooses at random among the **uphill moves**
 - The probability of picking a specific move can depend on the **steepness**
 - Converges **more slowly** than steepest ascent
 - Can find **higher quality** solutions
- **First-Choice Hill Climbing**
 - Implements stochastic hill climbing by randomly generating successors until a better successor than the current state is generated
 - It performs **well** when a state has many successors
- **Random-Restart Hill Climbing**
 - Generates a series of hill-climbing searches from random initial states
 - It stops when a goal is found
 - It is **complete** with probability 1 as it will eventually generate a goal state as initial state

OUTLINE

- 1 Recap: Informed Search and Optimisation
- 2 Hill Climbing: Variants
- 3 Simulated Annealing**
- 4 Constraint Handling

SIMULATED ANNEALING

- A hill climbing algorithm that **never makes downhill moves** can get stuck in local maxima (or minima)
- A purely random walk will eventually find the global maximum, but will be **extremely inefficient**
- **Simulated Annealing** combines these two approaches:
 - Generate initial candidate solution at random
 - Generate neighbour solutions, **pick one at random (instead of best)**
 - If the move improves the current solution, it is always accepted
 - Otherwise, the algorithm accepts a worse move with **a probability less than 1 that decreases over time**
 - Terminate when **a min temperature is reached** or **solution does not change in value**
- Differently from hill climbing, simulated annealing allows bad moves

SIMULATED ANNEALING: ACCEPTING BAD MOVES

- How should we set the probability of accepting bad moves?
 - **High** probability in the beginning, similar to random search (explore)
 - **Lower** as time passes, similar to hill climbing (exploit)
- In metallurgy, **annealing** is the process used to temper or harden metals by heating them to a **high** temperature
- Then, **gradually** cooling them until they reach a low-energy state
- Simulated annealing uses a similar principle to accept worse moves
 - The algorithm accepts worse moves with some probability less than 1
 - The probability **decreases exponentially** the worse the move is
 - The probability decreases as the **temperature goes down**

PROBABILITY FUNCTION: QUALITY

- We model the probability using a thermodynamics-inspired function:

$$e^{\Delta E/T}$$

- 'The probability **decreases exponentially** the worse the move is'. This is determined by the numerator in the exponential:

$$\Delta E = E(x_{\text{new}}) - E(x_{\text{current}}) \leq 0,$$

where $E(x_{\text{new}})$ is the energy (quality) of the new random state and $E(x_{\text{current}})$ is the energy (quality) of the current state

- **Case 1: Very bad move.** If $E(x_{\text{new}}) \ll E(x_{\text{current}})$, then $e^{\Delta E/T} \approx 0$, i.e., a lower probability to accept the move
- **Case 2: Not too bad move.** If $|E(x_{\text{new}}) - E(x_{\text{current}})| \approx 0$, then $e^{\Delta E/T} > 0$, i.e., a higher probability to accept the move

PROBABILITY FUNCTION: TEMPERATURE

- 'The probability decreases as the **temperature goes down**'. This is determined by the denominator $T > 0$
- At start, we set the temperature to a high value $T(0) = \mathcal{T} \gg 0$
- An update rule (**schedule**) is used to reduce it, for instance:

$$T(t+1) = \alpha T(t),$$

where $T(t+1)$ and $T(t)$ represent the values of the temperature at the next iteration and at this iteration, respectively, and $\alpha \approx 1$

- If the temperature decreases slowly enough, simulated annealing will **find the global maximum with probability 1**
- This is a property of the Boltzman distribution $e^{\Delta E/T}$, i.e., the probability is **concentrated on the global maxima**

PROBABILITY FUNCTION: TEMPERATURE

- We model the probability using a thermodynamics-inspired function:

$$e^{\Delta E/T}$$

- How does the temperature affect the probability?
- **Case 1: High temperature.** If $T \gg 0$, then $e^{\Delta E/T} \approx 1$, i.e., a higher probability to accept the move
- **Case 2: Low temperature.** If $T \approx 0$, then $e^{\Delta E/T} \approx 0$, i.e., a lower probability to accept the move

SIMULATED ANNEALING

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state
 $current \leftarrow problem.INITIAL$
 for $t = 1$ **to** ∞ **do**
 $T \leftarrow schedule(t)$
 if $T = 0$ **then return** $current$
 $next \leftarrow$ a randomly selected successor of $current$
 $\Delta E \leftarrow VALUE(current) - VALUE(next)$
 if $\Delta E > 0$ **then** $current \leftarrow next$
 else $current \leftarrow next$ only with probability $e^{\Delta E/T}$

PERFORMANCE OF SIMULATED ANNEALING

- Let us evaluate the performance of simulated annealing
 - **Completeness:** simulated annealing **is not complete**, as it depends on the problem formulation and design of the algorithm
 - **Optimality:** simulated annealing **is not optimal**, as it depends on the termination criteria and the schedule
 - **Time complexity:** it depends on the schedule and the termination criterion
 - **Space complexity:** depends on how the design variable is represented in the algorithm

SIMULATED ANNEALING

- Pros:
 - Can usually find a **good (near optimal) solution** in a reasonable amount of time
 - It avoids getting stuck **in poor local maxima and plateau** by combining exploration and exploitation
- Cons:
 - **Not guaranteed** to be complete, nor optimal
 - Time/space complexity is **problem- and representation-dependent**

OUTLINE

- 1 Recap: Informed Search and Optimisation
- 2 Hill Climbing: Variants
- 3 Simulated Annealing
- 4 Constraint Handling**

APPLYING OPTIMISATION ALGORITHMS

- To apply hill climbing and simulated annealing we need to specify:
- **Optimisation problem formulation**
 - Design variable and search space
 - Objective function
 - Constraints
- **Algorithm operators**
 - Representation
 - Initialisation procedure
 - Neighbourhood operator
- **Strategy to deal with constraints**
 - 1 By carefully designing algorithm operators
 - 2 By modifying the objective function

OPTIMISATION: AN EXAMPLE

Activity. Find the closest point to the origin that lies on the line $y = 2x - 4$. In pairs or small groups, formulate the optimisation problem.

- **Design variable**

$$\bar{x} \in \mathbb{R}^2, \bar{x} = [x, y]^\top$$

- **Search space**

$$\mathbb{R}^2$$


- **Objective function**

$$\min f(\bar{x}) = \sqrt{(x - x_0)^2 + (y - y_0)^2}$$

- **Constraints**

$$h(\bar{x}) = 0, \text{ where}$$

$$h(\bar{x}) = \begin{cases} 0, & \text{if } y = 2x - 4, \\ 1, & \text{otherwise.} \end{cases}$$



OPTIMISATION: AN EXAMPLE

Activity. Find the closest point to the origin that lies on the line $y = 2x - 4$. In pairs or small groups, formulate the optimisation problem without any explicit constraints.

- Design variable**

$$\bar{x} \in \{x, y : y = 2x - 4\}, \bar{x} = [x, y]^T$$

- Search space**


$$\{x, y : y = 2x - 4\}$$

- Objective function**

$$\min f(\bar{x}) = \sqrt{(x - x_0)^2 + (y - y_0)^2}$$

- Constraints**

No explicit constraints



OPTIMISATION: AN EXAMPLE

Example. Find the closest point to the origin that lies on the line $y = 2x - 4$. Let us determine the algorithm operators.

- **Representation**

Real vector of size 2


- **Initialisation procedure**

Initialise with two real values picked uniformly at random

- **Neighbourhood operator**

Add or subtract 0.1 to either variable

- Hill climbing and simulated annealing
may return an infeasible solution



CONSTRAINT HANDLING: ALGORITHM OPERATORS

- **Representation**

Real vector of size 2

- **Initialisation procedure**

Initialise with two real values

$$\bar{x} \in \{x, y : y = 2x - 4\}$$


picked uniformly at random

- **Neighbourhood operator**

Add or subtract 0.1 to x
and find corresponding y

- **Constraint strategy**

Operators above



CONSTRAINT HANDLING: ALGORITHM OPERATORS

- Pros:
 - **Infeasible** candidate solutions **will not be generated**, facilitating the search for optimal solutions
- Cons:
 - **May be difficult** to design, **problem-dependent**
 - Can **restrict the search space too much**, making it harder to find the optimal solution

CONSTRAINT HANDLING: OBJECTIVE FUNCTION

- Design variable**

$$\bar{x} \in \mathbb{R}^2, \bar{x} = [x, y]^\top$$

- Search space**

$$\mathbb{R}^2$$


- Objective function**

$$\min f(\bar{x}) = \sqrt{(x - x_0)^2 + (y - y_0)^2}$$

- Constraints**

$$h(\bar{x}) = 0, \text{ where}$$

$$h(\bar{x}) = \begin{cases} 0, & \text{if } y = 2x - 4, \\ 1, & \text{otherwise.} \end{cases}$$



CONSTRAINT HANDLING: OBJECTIVE FUNCTION

- **Representation**

Real vector of size 2


- **Initialisation procedure**

Initialise with two real values
picked uniformly at random

- **Neighbourhood operator**

Add or subtract 0.1 to either variable

- How do we avoid infeasible solutions?



APPROACH 1: DEATH PENALTY

- Given the canonical formulation of an optimisation problem:

$$\begin{aligned} \min \quad & f(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n. \end{aligned}$$

- We avoid infeasible solutions by modifying the objective function

$$\min \quad f(x) + Q(x),$$

where the penalty term $Q(x)$ is defined as:

$$Q(x) \begin{cases} 0, & \text{if } x \text{ is feasible,} \\ C, & \text{otherwise.} \end{cases}$$

- The term C is a large positive constant, ensuring infeasible solutions are worse than feasible solutions

APPROACH 1: DEATH PENALTY

- Given the canonical formulation of an optimisation problem:

$$\begin{aligned} \min \quad & f'(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n. \end{aligned}$$

- We can also use the following formulation:

$$f'(x) = \begin{cases} f(x), & \text{if } x \text{ is feasible,} \\ C, & \text{otherwise.} \end{cases}$$

- The term C is a large positive constant, like before
- Weakness: all infeasible solutions have the **same penalty**
- Algorithms **will struggle** to find feasible solutions in regions of the state space dominated by infeasible solutions

APPROACH 2: LEVELS OF INFEASIBILITY

- Given the canonical formulation of an optimisation problem:

$$\begin{aligned} \min \quad & f(x) + Q(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n. \end{aligned}$$

- The idea is to distinguish the objective value of infeasible solutions to help the algorithm find a feasible solution

$$\begin{aligned} Q(x) = & C g_1(x) + C g_2(x) + \dots + C g_m(x) \\ & + C |h_1(x)| + C |h_2(x)| + \dots + C |h_n(x)|. \end{aligned}$$

- The **objective function itself** will guide the algorithm towards feasible solutions (and near-optimal solutions, too!)

APPROACH 2: LEVELS OF INFEASIBILITY

- Given the canonical formulation of an optimisation problem:

$$\begin{aligned} \min \quad & f(x) + Q(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n. \end{aligned}$$

- We should only add the penalties corresponding to violated constraints

$$\begin{aligned} Q(x) = & v_{g_1} C |g_1(x)| + v_{g_2} C |g_2(x)| + \dots + v_{g_m} C |g_m(x)| \\ & + v_{h_1} C |h_1(x)| + v_{h_2} C |h_2(x)| + \dots + v_{h_n} C |h_n(x)|, \end{aligned}$$

where $v_{g_i} = 1$ if g_i is violated and $v_{g_i} = 0$ otherwise,
and $v_{h_j} = 1$ if h_j is violated and $v_{h_j} = 0$ otherwise

APPROACH 2: LEVELS OF INFEASIBILITY

- Given the canonical formulation of an optimisation problem:

$$\begin{aligned} \min \quad & f(x) + Q(x), \\ \text{s.t.} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m, \\ & h_j(x) = 0, \quad j = 1, \dots, n. \end{aligned}$$

- We should only add the penalties corresponding to violated constraints

$$\begin{aligned} Q(x) = & v_{g_1} C g_1(x)^2 + v_{g_2} C g_2(x)^2 + \dots + v_{g_m} C g_m(x)^2 \\ & + v_{h_1} C h_1(x)^2 + v_{h_2} C h_2(x)^2 + \dots + v_{h_n} C h_n(x)^2. \end{aligned}$$

- Squared values** make the distinction between objective values even larger, distinguishing different infeasible solutions more

CONSTRAINT HANDLING: OBJECTIVE FUNCTION

- Pros:
 - Usually easier to design
- Cons:
 - Algorithms have to search for feasible solutions to design

FINAL REMARKS: COMPLETENESS

- If we use a strategy that never enables infeasible solutions to be generated, then hill climbing and simulated annealing are **complete**
- Otherwise:
 - Hill climbing **is not complete** if the objective function has local optima (a local optimum may be infeasible)
 - Simulated annealing **is not guaranteed** to find a feasible solution within a reasonable amount of time; it might terminate before finding a feasible solution

References



Russell, A. S., and Norvig, P., *Artificial Intelligence A Modern Approach*, 4th Edition. Prentice Hall.



Chapter 4 – “Search in Complex Environments”, Section 4.1 “Local Search and Optimization Problems” up to and including Section 4.1.2 “Simulated Annealing”.

AIMS OF THE SESSION

You should now be able to:

- Understand the weaknesses in Hill Climbing and how to mitigate them
- Explain the steps involved in Simulated Annealing
- Deal with constraints in optimisation problems



Thank you!